



# INSTALLATION GUIDE

A guide for installing or migrating to CircleCI  
Server v4.0.0 on Amazon Web Services

[docs@circleci.com](mailto:docs@circleci.com)

Version 4.0.0, 08/18/2022: FINAL

<b>Phase 1 - Prerequisites</b>	<b>1</b>
1. Install required software	1
AWS required software	1
S3 compatible storage required software	1
2. Create a Kubernetes cluster	1
EKS	2
Enable Workload Identity in GKE (optional)	2
3. Create a new GitHub OAuth app	3
GitHub Enterprise	5
4. Frontend TLS certificates	5
AWS Route 53	6
5. Encryption/signing keys	6
a. Artifact signing key	6
b. Encryption signing key	6
6. Object storage and permissions	6
S3 Storage	7
a. Create AWS S3 Bucket	7
b. Set up authentication	7
<b>Phase 2 - Core services</b>	<b>15</b>
1. Create a namespace	15
2. Pull images	15
3. Create helm values	16
a. API token	16
b. Session cookie	17
c. Encryption	17
d. Postgres credentials	17
e. MongoDB credentials	18
f. RabbitMQ configurations and auth Secrets	19
g. Pusher Kubernetes Secret	19
h. Global	20
CircleCI domain name (required)	20
License	20
Registry	20
i. TLS	21
j. GitHub integration	22
GitHub	22
GitHub Enterprise	22
k. Object storage	23
S3 compatible	23
l. Installing behind a proxy	25
4. Deploy	26
5. Create DNS entry	26

6. Validation .....	26
<b>Phase 3 - Execution environments .....</b>	<b>27</b>
1. Nomad clients .....	27
a. Create your cluster with Terraform .....	27
AWS cluster .....	27
b. Nomad Autoscaler configuration .....	28
AWS autoscaler IAM/role .....	28
2. Nomad Servers .....	30
a. Nomad gossip encryption key .....	30
b. Nomad mTLS .....	31
c. Nomad Autoscaler .....	31
AWS .....	31
d. Helm upgrade .....	32
3. Nomad Clients Validation .....	32
GitHub Cloud .....	33
GitHub Enterprise .....	33
3. VM service .....	34
AWS .....	34
Set up security group .....	34
Set up authentication .....	36
VM Service Validation .....	44
4. Runner .....	44
Overview .....	44
<b>Phase 4 - Post installation .....</b>	<b>45</b>
Backup and restore .....	45
Set up backup and restore on AWS .....	45
Set up backup and restore with S3-compatible storage .....	48
Take a backup .....	50
Email Notifications .....	51
Managing orbs .....	51
<b>Hardening your cluster .....</b>	<b>52</b>
Network topology .....	52
Network traffic .....	52
Reverse proxy status .....	53
Kubernetes load balancers .....	53
Ingress .....	53
Egress .....	53
Common rules for compute instances .....	53
Ingress .....	54
Egress .....	54
Kubernetes nodes .....	54
Intra-node traffic .....	54

Ingress .....	54
Egress .....	54
Nomad Clients .....	55
Ingress .....	55
Egress .....	55
External VMs .....	55
Ingress .....	55
Egress .....	55
<b>Installing server behind a proxy .....</b>	<b>56</b>
Known limitations .....	56
<b>Migrate from server v3.x to v4.x .....</b>	<b>58</b>
Prerequisites .....	58
Migration .....	58
1. Create docker secret for CircleCI image registry .....	59
2. Clone the repository and run the kots-exporter script .....	59
3. Validate your helm-value file .....	60
4. Generate helm-diff output .....	61
5. Upgrading CircleCI Server 3.x .....	61
6. Check upgrade status .....	62
7. Update DNS setting .....	62
8. Execute Nomad Terraform .....	62
9. Validate your migration to server v4.x .....	63
10. Update your team .....	63

# Phase 1 - Prerequisites

CircleCI server v4.x is installed as a helm chart. The installation process is broken down into four phases. There is a validation step at the end of each phase, allowing you to confirm success before moving to the next phase. Depending on your requirements, phases 3 and 4 may include multiple steps. This installation guide assumes you have already read the [CircleCI Server v4.x Overview](#).



In the following sections, replace any sections indicated by `< >` with your details.

## 1. Install required software

Download and install the following software before continuing:

Tool	Version	Used for	Notes
<a href="#">Terraform</a>	0.15.4 or greater	Infrastructure Management	
<a href="#">kubect</a>	1.19 or greater	Kubernetes CLI	
<a href="#">Helm</a>	3.9.2 or greater	Kubernetes Package Management	
<a href="#">Helm Diff</a>	3.5.0 or greater	Helping with <code>values.yaml</code> changes and updates	Optional, but may help with troubleshooting between releases
<a href="#">Velero CLI</a>	Latest	Backup and restore capability	See Velero's <a href="#">supported providers</a> documentation for further information.

### AWS required software

- [AWS CLI](#)

### S3 compatible storage required software

- Install and configure [MinIO CLI](#) for your storage provider.

## 2. Create a Kubernetes cluster

CircleCI server installs into an existing Kubernetes cluster. The application uses a large number of resources. Depending on your usage, your Kubernetes cluster should meet the following requirements:

Number of daily active CircleCI users	Minimum Nodes	Total CPU	Total RAM	NIC speed
< 500	3	12 cores	32 GB	1 Gbps
500+	3	48 cores	240 GB	10 Gbps

## Supported Kubernetes versions:

CircleCI Version	Kubernetes Version
4.0.0	1.22 - 1.23

Creating a Kubernetes cluster is your responsibility. Please note:

- You must have appropriate permissions to list, create, edit, and delete pods in your cluster. Run this command to verify your permissions:

```
kubectl auth can-i <list|create|edit|delete> pods
```

- There are no requirements regarding VPC setup or disk size for your cluster. It is recommended that you set up a new VPC rather than use an existing one.

## EKS

You can learn more about creating an Amazon EKS cluster in the [EKS docs](#). We recommend using `eksctl` to create your cluster, which creates a VPC and selects the proper security groups for you.

1. [Install](#) and [configure](#) the AWS CLI for your AWS account.
2. Install `eksctl`.
3. Create your cluster by running the following (Cloud formation with `eksctl` and EKS can take more than 20 minutes to complete):

```
eksctl create cluster --name=circleci-server --nodes 4 --node-type m5.xlarge
```

4. Once the cluster has been created, you can use the following command to configure `kubectl` access:

```
eksctl utils write-kubeconfig --cluster circleci-server
```



You may see the following error `AWS STS Access - cannot get role ARN for current session: InvalidClientTokenID`. This means your AWS credentials are invalid, or your IAM user does not have permission to create an EKS cluster. Proper IAM permissions are necessary to use `eksctl`. See the AWS documentation on [IAM permissions](#).

## Enable Workload Identity in GKE (optional)

Follow these steps if you already have a GKE cluster and need to enable Workload Identity on the cluster and the node pools.

1. Enable Workload Identity on existing cluster:

```
gcloud container clusters update "<CLUSTER_NAME>" \
  --workload-pool="<PROJECT_ID>.svc.id.goog"
```

2. Get node pools of existing GKE cluster:

```
gcloud container node-pools list --cluster "<CLUSTER_NAME>"
```

3. Update existing node pools:

```
gcloud container node-pools update "<NODEPOOL_NAME>" \
  --cluster="<CLUSTER_NAME>" \
  --workload-metadata="GKE_METADATA"
```

You must repeat Step 3 for all the existing node pools. Follow these links for steps to enable Workload Identity for your Kubernetes service accounts:

- [Nomad Autoscaler](#)
- [VM](#)
- [Object-Storage](#)

### 3. Create a new GitHub OAuth app



If GitHub Enterprise and CircleCI server are not on the same domain, then images and icons from GHE will fail to load in the CircleCI web app.

Registering and setting up a new GitHub OAuth app for CircleCI server allows for authorization control to your server installation using GitHub OAuth and for updates to GitHub projects/repos using build status information. The following steps apply for both GitHub.com and GitHub Enterprise.

1. In your browser, navigate to **your GitHub instance > User Settings > Developer Settings > OAuth Apps** and click the **New OAuth App** button.

## Register a new OAuth application

**Application name \***

circleci-server

Something users will recognize and trust.

**Homepage URL \***

https://circleci.someplace.com

The full URL to your application homepage.

**Application description**

Application description is optional

This is displayed to all users of your application.

**Authorization callback URL \***

https://circleci.someplace.com/auth/github

Your application's callback URL. Read our [OAuth documentation](#) for more information.

Register application


Cancel

Figure 1. New GitHub OAuth App

2. Complete the following fields, based on your planned installation:
  - **Homepage URL:** The URL of your planned CircleCI installation.
  - **Authorization callback URL:** The authorization callback URL is the URL of your planned CircleCI installation followed by `/auth/github`
3. Once completed, you will be shown the **Client ID**. Select **Generate a new Client Secret** to generate a Client Secret for your new OAuth App. You will need these values when you configure CircleCI server.



circleci-server

 **nathanfishcircleci** owns this application.

Transfer ownership

You can list your application in the [GitHub Marketplace](#) so that other users can discover it.

List this application in the Marketplace

0 users

Revoke all user tokens

Client ID

1d2249984d6ce6a5d599

Client secrets

Generate a new client secret

You need a client secret to authenticate as the application to the API.

Figure 2. Client ID and Secret

## GitHub Enterprise

If using GitHub Enterprise, you also need a personal access token and the domain name of your GitHub Enterprise instance.

Create the `defaultToken` by navigating to **User Settings > Developer Settings > Personal access tokens**. The default token requires no scopes. You will need this value when you configure CircleCI server.

## 4. Frontend TLS certificates

By default, CircleCI server creates self-signed certificates to get you started. In production, you should supply a certificate from a trusted certificate authority. The [Let's Encrypt](#) certificate authority, for example, can issue a free certificate using their [certbot](#) tool. The sections below cover using Google Cloud DNS and AWS Route 53.



It is important that your certificate contains both your domain and the `app.*` subdomain as subjects. For example, if you host your installation at `server.example.com`, your certificate must cover `app.server.example.com` and `server.example.com`.

Once you have created your certificates using one of the methods described below, you can use the following commands to retrieve the certificates later when you need them during this installation:

```
ls -l /etc/letsencrypt/live/<CIRCLECI_SERVER_DOMAIN>
```

```
cat /etc/letsencrypt/live/<CIRCLECI_SERVER_DOMAIN>/fullchain.pem
```

```
cat /etc/letsencrypt/live/<CIRCLECI_SERVER_DOMAIN>/privkey.pem
```

## AWS Route 53

1. If you are using AWS Route 53 for DNS, you need the **certbot-route53** plugin installed. You can install the plugin with the following command:

```
python3 -m pip install certbot-dns-route53
```

2. Then execute this example to create a private key and certificate (including intermediate certificates) locally in `/etc/letsencrypt/live/<CIRCLECI_SERVER_DOMAIN>`:

```
certbot certonly --dns-route53 -d "<CIRCLECI_SERVER_DOMAIN>" -d "app.<CIRCLECI_SERVER_DOMAIN>"
```

## 5. Encryption/signing keys

The keysets generated in this section are used to encrypt and sign artifacts generated by CircleCI. You will need these values to configure server.



Store these values securely. If they are lost, job history and artifacts will not be recoverable.

### a. Artifact signing key

To generate an artifact signing key, run the following command:

```
docker run circleci/server-keysets:latest generate signing -a stdout
```

### b. Encryption signing key

To generate an encryption signing key, run the following command:

```
docker run circleci/server-keysets:latest generate encryption -a stdout
```

## 6. Object storage and permissions

Server v4.x hosts build artifacts, test results, and other state object storage. The following storage options are supported:

- [AWS S3](#)
- [MinIO](#)
- [Google Cloud Storage](#)

While any S3 compatible object storage may work, we test and support AWS S3 and MinIO. For object storage providers that do not support the S3 API, such as Azure blob storage, we recommend using MinIO Gateway.

Follow the instructions below to create a bucket and access method for S3 or GCS.



If you are installing behind a proxy, object storage should be behind this proxy also. Otherwise, proxy details will need to be supplied at the job level within every project `.circleci/config.yml` to allow artifacts, test results, cache save and restore, and workspaces to work. For more information see the [Installing Server Behind a Proxy](#) guide.

## S3 Storage

### a. Create AWS S3 Bucket

```
aws s3api create-bucket \
  --bucket <YOUR_BUCKET_NAME> \
  --region <YOUR_REGION> \
  --create-bucket-configuration LocationConstraint=<YOUR_REGION>
```

### b. Set up authentication

There are two ways to authenticate CircleCI with S3: IAM Roles for Service Accounts (IRSA), and IAM access keys. IRSA is the recommended method.

#### Option 1: IRSA

The following is a summary of [AWS's Documentation on IRSA](#) that is sufficient for installing CircleCI.

1. Create an IAM OIDC identity provider for your EKS Cluster.

```
eksctl utils associate-iam-oidc-provider --cluster <CLUSTER_NAME> --approve
```

2. Get the OIDC provider ARN. You will need this in later steps.

```
aws iam list-open-id-connect-providers | grep $(aws eks describe-cluster --name <CLUSTER_NAME> \
  --query "cluster.identity.oidc.issuer" --output text | awk -F '/' '{print $NF}')
```

3. Get your OIDC provider URL. You will need this in later steps.

```
aws eks describe-cluster --name <CLUSTER_NAME> --query "cluster.identity.oidc.issuer" --output
text | sed -e "s/^https:\/\/\/"
```

4. Create the role using the command and trust policy template below. You will need the Role ARN and name in later steps.

```
aws iam create-role --role-name circleci-s3 --assume-role-policy-document
file:///<TRUST_POLICY_FILE>
```

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Federated": "<OIDC_PROVIDER_ARN>"
      },
      "Action": "sts:AssumeRoleWithWebIdentity",
      "Condition": {
        "StringEquals": {
          "<OIDC_PROVIDER_URL>:sub": "system:serviceaccount:<K8S_NAMESPACE>:object-storage"
        }
      }
    }
  ]
}
```

5. Create the policy using the command and template below. Fill in the bucket name and the role ARN.

```
aws iam create-policy --policy-name circleci-s3 --policy-document file:///<POLICY_FILE>
```

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:PutAnalyticsConfiguration",
        "s3:GetObjectVersionTagging",
        "s3:CreateBucket",
        "s3:GetObjectAcl",
        "s3:GetBucketObjectLockConfiguration",
```

```
"s3:DeleteBucketWebsite",
"s3:PutLifecycleConfiguration",
"s3:GetObjectVersionAcl",
"s3:PutObjectTagging",
"s3:DeleteObject",
"s3:DeleteObjectTagging",
"s3:GetBucketPolicyStatus",
"s3:GetObjectRetention",
"s3:GetBucketWebsite",
"s3:GetJobTagging",
"s3:DeleteObjectVersionTagging",
"s3:PutObjectLegalHold",
"s3:GetObjectLegalHold",
"s3:GetBucketNotification",
"s3:PutBucketCORS",
"s3:GetReplicationConfiguration",
"s3:ListMultipartUploadParts",
"s3:PutObject",
"s3:GetObject",
"s3:PutBucketNotification",
"s3:DescribeJob",
"s3:PutBucketLogging",
"s3:GetAnalyticsConfiguration",
"s3:PutBucketObjectLockConfiguration",
"s3:GetObjectVersionForReplication",
"s3:GetLifecycleConfiguration",
"s3:GetInventoryConfiguration",
"s3:GetBucketTagging",
"s3:PutAccelerateConfiguration",
"s3:DeleteObjectVersion",
"s3:GetBucketLogging",
"s3:ListBucketVersions",
"s3:ReplicateTags",
"s3:RestoreObject",
"s3:ListBucket",
"s3:GetAccelerateConfiguration",
"s3:GetBucketPolicy",
"s3:PutEncryptionConfiguration",
"s3:GetEncryptionConfiguration",
"s3:GetObjectVersionTorrent",
"s3:AbortMultipartUpload",
```

```

        "s3:PutBucketTagging",
        "s3:GetBucketRequestPayment",
        "s3:GetAccessPointPolicyStatus",
        "s3:GetObjectTagging",
        "s3:GetMetricsConfiguration",
        "s3:PutBucketVersioning",
        "s3:GetBucketPublicAccessBlock",
        "s3:ListBucketMultipartUploads",
        "s3:PutMetricsConfiguration",
        "s3:PutObjectVersionTagging",
        "s3:GetBucketVersioning",
        "s3:GetBucketAcl",
        "s3:PutInventoryConfiguration",
        "s3:GetObjectTorrent",
        "s3:PutBucketWebsite",
        "s3:PutBucketRequestPayment",
        "s3:PutObjectRetention",
        "s3:GetBucketCORS",
        "s3:GetBucketLocation",
        "s3:GetAccessPointPolicy",
        "s3:GetObjectVersion",
        "s3:GetAccessPoint",
        "s3:GetAccountPublicAccessBlock",
        "s3:ListAllMyBuckets",
        "s3:ListAccessPoints",
        "s3:ListJobs"
    ],
    "Resource": [
        "arn:aws:s3:::<YOUR_BUCKET_NAME>",
        "arn:aws:s3:::<YOUR_BUCKET_NAME>/*"
    ]
},
{
    "Effect": "Allow",
    "Action": [
        "iam:GetRole",
        "sts:AssumeRole"
    ],
    "Resource": "<OBJECT_STORAGE_ROLE_ARN>"
}
]

```

```
}
```

6. Attach the policy to the role:

```
aws iam attach-role-policy --role-name <OBJECT_STORAGE_ROLE_NAME> --policy-arn  
=<STORAGE_POLICY_ARN>
```

## Option 2: IAM access keys

1. Create an IAM user for CircleCI server.

```
aws iam create-user --user-name circleci-server
```

2. Create a policy document `policy.json`.

```
{  
  "Version": "2012-10-17",  
  "Statement": [  
    {  
      "Effect": "Allow",  
      "Action": [  
        "s3:PutAnalyticsConfiguration",  
        "s3:GetObjectVersionTagging",  
        "s3:CreateBucket",  
        "s3:GetObjectAcl",  
        "s3:GetBucketObjectLockConfiguration",  
        "s3>DeleteBucketWebsite",  
        "s3:PutLifecycleConfiguration",  
        "s3:GetObjectVersionAcl",  
        "s3:PutObjectTagging",  
        "s3>DeleteObject",  
        "s3>DeleteObjectTagging",  
        "s3:GetBucketPolicyStatus",  
        "s3:GetObjectRetention",  
        "s3:GetBucketWebsite",  
        "s3:GetJobTagging",  
        "s3>DeleteObjectVersionTagging",  
        "s3:PutObjectLegalHold",  
        "s3:GetObjectLegalHold",  
        "s3:GetBucketNotification",  
      ]  
    }  
  ]  
}
```

```
"s3:PutBucketCORS",
"s3:GetReplicationConfiguration",
"s3:ListMultipartUploadParts",
"s3:PutObject",
"s3:GetObject",
"s3:PutBucketNotification",
"s3:DescribeJob",
"s3:PutBucketLogging",
"s3:GetAnalyticsConfiguration",
"s3:PutBucketObjectLockConfiguration",
"s3:GetObjectVersionForReplication",
"s3:GetLifecycleConfiguration",
"s3:GetInventoryConfiguration",
"s3:GetBucketTagging",
"s3:PutAccelerateConfiguration",
"s3>DeleteObjectVersion",
"s3:GetBucketLogging",
"s3:ListBucketVersions",
"s3:ReplicateTags",
"s3:RestoreObject",
"s3:ListBucket",
"s3:GetAccelerateConfiguration",
"s3:GetBucketPolicy",
"s3:PutEncryptionConfiguration",
"s3:GetEncryptionConfiguration",
"s3:GetObjectVersionTorrent",
"s3:AbortMultipartUpload",
"s3:PutBucketTagging",
"s3:GetBucketRequestPayment",
"s3:GetAccessPointPolicyStatus",
"s3:GetObjectTagging",
"s3:GetMetricsConfiguration",
"s3:PutBucketVersioning",
"s3:GetBucketPublicAccessBlock",
"s3:ListBucketMultipartUploads",
"s3:PutMetricsConfiguration",
"s3:PutObjectVersionTagging",
"s3:GetBucketVersioning",
"s3:GetBucketAcl",
"s3:PutInventoryConfiguration",
"s3:GetObjectTorrent",
```



```

        "s3:PutBucketWebsite",
        "s3:PutBucketRequestPayment",
        "s3:PutObjectRetention",
        "s3:GetBucketCORS",
        "s3:GetBucketLocation",
        "s3:GetAccessPointPolicy",
        "s3:GetObjectVersion",
        "s3:GetAccessPoint",
        "s3:GetAccountPublicAccessBlock",
        "s3:ListAllMyBuckets",
        "s3:ListAccessPoints",
        "s3:ListJobs"
    ],
    "Resource": [
        "arn:aws:s3:::<YOUR_BUCKET_NAME>",
        "arn:aws:s3:::<YOUR_BUCKET_NAME>/*"
    ]
}
]
}

```

3. Attach policy to user.

```

aws iam put-user-policy \
  --user-name circleci-server \
  --policy-name circleci-server \
  --policy-document file://policy.json

```

4. Create Access Key for user circleci-server.



You will need this when you configure your server installation later.

```

aws iam create-access-key --user-name circleci-server

```

The result should look like this:

```
{
  "AccessKey": {
    "UserName": "circleci-server",
    "Status": "Active",
    "CreateDate": "2017-07-31T22:24:41.576Z",
    "SecretAccessKey": <AWS_SECRET_ACCESS_KEY>,
    "AccessKeyId": <AWS_ACCESS_KEY_ID>
  }
}
```

# Phase 2 - Core services

Before you begin with the CircleCI server v4.x core services installation phase, ensure all [prerequisites](#) are met.



In the following sections, replace any sections indicated by `< >` with your details.

## 1. Create a namespace

Create a namespace to install the application into.

```
kubectl create ns <namespace>
```



Once you have created your namespace, we recommend setting your `kubectl` context too, with the following command: `kubectl config set-context --current --namespace <namespace>`

## 2. Pull images

Credentials to pull the images from CircleCI's image registry will be provided to you as part of the onboarding process. A `docker-registry` Kubernetes Secret will be used to pull images from Azure Container Registry (ACR). You have two options, depending on whether your application has access to the public internet.

**Option 1:** Your application has access to the public internet.

This example creates a Kubernetes Secret to enable deployments to pull images from CircleCI's image registry. The `docker-registry` Kubernetes Secret takes the following form:

```
kubectl create secret docker-registry regcred \
  --docker-server=cciserver.azurecr.io \
  --docker-username=<your-username> \
  --docker-password="<provided-token>" \
  --docker-email=<your-contact-email>
```

**Option 2:** Your application does NOT have access to the public internet.

The credentials provided to you allow you to pull and store copies of our images locally. Pull and store the images in whichever Docker repository you have available. The `docker-registry` Kubernetes Secret takes the following form:

```
kubectl create secret docker-registry regcred \
  --docker-server=<your-docker-image-repo> \
  --docker-username=<your-username> \
  --docker-password=<your-access-token> \
  --docker-email=<your-email>
```

### 3. Create helm values

Before installing CircleCI, it is recommended to create a new `values.yaml` file unique to your installation. [The Installation Reference section](#) contains some example `values.yaml` files that are a good place to start. The following describes the minimum required values to include in `values.yaml`. Additional customizations are available, see the provided `values.yaml` for all available options.

For sensitive data there are two options:

- add into the `values.yaml` file
- add them as Kubernetes Secrets directly

This flexibility allows you to manage Kubernetes Secrets using whichever process you prefer.



Whichever option you choose, this sensitive information is stored as a Kubernetes Secret within CircleCI.

#### a. API token

The application requires a Kubernetes Secret containing an API token. This API token is used to facilitate internal API communication to `api-service`. Use a random string and store it securely, CircleCI will not be able to recover this value if lost. There are two options depending on whether you want to create the Kubernetes Secret, or if you want CircleCI to create it for you.

**Option 1:** Create the Kubernetes Secret yourself.

```
kubectl create secret generic api-token \
  --from-literal=api-token=<your-super-secret-random-value>
```

**Option 2:** CircleCI creates the Kubernetes Secret for you.

Add the value to `values.yaml`. CircleCI will create the Kubernetes Secret automatically.

```
apiToken: "<your-super-secret-random-value>"
```

## b. Session cookie

The application requires a session cookie key Kubernetes Secret, which CircleCI uses to sign session cookies. The Secret must be exactly 16 characters long. Use a random string and store it securely, CircleCI will not be able to recover this value if lost. There are two options depending on whether you want to create the Kubernetes Secret, or if you want CircleCI to create it for you.

**Option 1:** Create the Kubernetes Secret yourself.

```
kubectl create secret generic session-cookie \
  --from-literal=session-cookie-key=<your-secret-key-16-chars>
```

**Option 2:** CircleCI creates the Kubernetes Secret for you.

Add the value to `values.yaml`. CircleCI will create the Secret automatically.

```
sessionCookieKey: "<your-secret-key-16-chars>"
```

## c. Encryption

The application requires a Kubernetes Secret containing signing and encryption keysets. These keysets are used to encrypt and sign artifacts generated by CircleCI. These keys were created during the [prerequisites phase](#). CircleCI will not be able to recover the values if lost. Depending on how you prefer to manage Kubernetes Secrets, there are two options.

**Option 1:** Create the Kubernetes Secret yourself.

```
kubectl create secret generic signing-keys \
  --from-literal=signing-key=<your-generated-signing-key> \
  --from-literal=encryption-key=<your-generated-encryption-key>
```

**Option 2:** CircleCI creates the Kubernetes Secret.

Add the value to `values.yaml`. CircleCI will create the Secret automatically.

```
keyset:
  signing: '<your-generated-signing-key>'
  encryption: '<your-generated-encryption-key>'
```

## d. Postgres credentials

The application requires a Kubernetes Secret containing Postgres credentials. This is true when using either the internal (default) or an externally hosted instance of Postgres. CircleCI will not be able to recover the values if lost. Based on how you prefer to manage Kubernetes Secrets there are two options.

**Option 1:** Create the Secret yourself.

```
kubectl create secret generic postgresql \
  --from-literal=postgres-password=<postgres-password>
```

You must then provide the following to the `values.yaml` file:

```
postgresql:
  auth:
    existingSecret: postgresql
```

**Option 2:** CircleCI creates the Kubernetes Secret.

Add the credentials to `values.yaml`, and CircleCI will create the Secret automatically.

```
postgresql:
  auth:
    postgresPassword: "<postgres-password>"
```

## e. MongoDB credentials

The application requires a Kubernetes Secret containing MongoDB credentials. This is true when using either the internal (default) or an externally hosted instance of MongoDB. CircleCI will not be able to recover the values if lost. Based on how you prefer to manage Kubernetes Secrets there are two options.

**Option 1:** Create the Kubernetes Secret yourself.

```
kubectl create secret generic mongodb-credentials \
  --from-literal=mongodb-root-password=<root-password> \
  --from-literal=mongodb-password=<user-password>
```

You must then provide the following to the `values.yaml` file:

```
mongodb:
  auth:
    existingSecret: mongodb-credentials
```

**Option 2:** CircleCI creates the Kubernetes Secret.

Add the credentials to `values.yaml`, and CircleCI will create the Secret automatically.

```
mongodb:
  auth:
    rootPassword: "<root-password>"
    password: "<user-password>"
```

## f. RabbitMQ configurations and auth Secrets

The RabbitMQ installation requires two random alphanumeric strings. CircleCI will not be able to recover the values if lost. Based on how you prefer to manage Kubernetes Secrets there are two options.

**Option 1:** Create the Secret yourself.

```
kubectl create secret generic rabbitmq-key \
--from-literal=rabbitmq-password=<secret-alphanumeric-password> \
--from-literal=rabbitmq-erlang-cookie=<secret-alphanumeric-key>
```

You must then provide the following to the `values.yaml` file:

```
rabbitmq:
  auth:
    existingPasswordSecret: rabbitmq-key
    existingErlangSecret: rabbitmq-key
```

**Option 2:** CircleCI creates the Kubernetes Secret.

Add the value to `values.yaml`, and CircleCI will create the Kubernetes Secret automatically.

```
rabbitmq:
  auth:
    password: "<secret-alphanumeric-password>"
    erlangCookie: "<secret-alphanumeric-key>"
```

## g. Pusher Kubernetes Secret

The application requires a Kubernetes Secret for Pusher. CircleCI will not be able to recover the values if lost. Based on how you prefer to manage Kubernetes Secrets there are 2 options:

**Option 1:** Create the Kubernetes Secret yourself.

```
kubectl create secret generic pusher \
--from-literal=secret=<pusher-secret>
```

**Option 2:** CircleCI creates the Kubernetes Secret.

Add the value to `values.yaml`, and CircleCI will create the Kubernetes Secret automatically.

```
pusher:
  secret: "<pusher-secret>"
```

## h. Global

All values in this section are children of `global` in your `values.yaml`.

### CircleCI domain name (required)

Enter the domain name you specified when creating your [Frontend TLS key and certificate](#).

```
global:
  ...
  domainName: "<full-domain-name-of-your-install>"
```

### License

A license will be provided by CircleCI, add it to `values.yaml`:

```
global:
  ...
  license: '<license>'
```

### Registry

The registry to pull images from will have been provided to you, or you may have added the images to your own hosted registry. Add the registry to `values.yaml`:

```
global:
  ...
  container:
    registry: <registry-domain eg: cciserver.azurecr.io >
    org: <your-org-if-applicable>
```



## i. TLS

For TLS, you have 4 options:

### Do nothing

Do nothing. [Let's Encrypt](#) will automatically request and manage certificates for you. This is a good option for trials but not recommended for production use.

### Supply a private key and certificate

You can supply a private key and certificate, which you may have created during the prerequisites steps. The key and certificates will need to be base64 encoded. You can retrieve and encode the values with the following commands:

```
cat /etc/letsencrypt/live/<CIRCLECI_SERVER_DOMAIN>/privkey.pem | base64
cat /etc/letsencrypt/live/<CIRCLECI_SERVER_DOMAIN>/fullchain.pem | base64
```

And add them to `values.yaml`:

```
tls:
  certificate: '<full-chain>'
  privateKey: '<private-key>'
```

### Use ACM

Have [AWS Certificate Manager \(ACM\)](#) automatically request and manage certificates for you. Follow the [ACM documentation](#) for instructions on how to generate ACM certificates.

Enable `aws_acm` and add the `service.beta.kubernetes.io/aws-load-balancer-ssl-cert` annotation to point at the ACM ARN:

```
nginx:
  annotations:
    service.beta.kubernetes.io/aws-load-balancer-ssl-cert: <acm-arn>
  aws_acm:
    enabled: false
```



If you have already deployed CircleCI server, enabling ACM is a destructive change to the loadbalancer. The service will have to be regenerated to allow the use of your ACM certificates and so the associated loadbalancer will also be regenerated. You will need to update your DNS records to the new loadbalancer once you have redeployed CircleCI server.

### Disable TLS within CircleCI

You can choose to disable TLS termination within CircleCI. The system will still need to be accessed over HTTPS, so TLS termination will be required somewhere upstream of CircleCI. Implement this by following the first option (do nothing) and forward the following ports to your CircleCI load balancer:

- Frontend / API Gateway [TCP 80, 443]
- VM service [TCP 3000]
- Nomad server [TCP 4647]
- Output processor [gRPC 8585]

## j. GitHub integration

To configure GitHub with CircleCI, there are two options for providing credentials to the deployment. Steps for both GitHub and GitHub Enterprise (GHE) are given in the next two sections.

### GitHub

These instructions are for the GitHub.com, **not** GitHub Enterprise. Use the client ID and secret you created with your Github OAuth application in the [prerequisites phase](#).

**Option 1:** Create the Kubernetes Secret yourself.

```
kubectl create secret generic github-secret \
  --from-literal=clientId=<client-id> \
  --from-literal=clientSecret=<client-secret>
```

**Option 2:** CircleCI creates the Kubernetes Secret.

Add the client ID and secret to the `values.yaml` file. CircleCI will create the Kubernetes Secret automatically.

```
github:
  clientId: "<client-id>"
  clientSecret: "<client-secret>"
```

### GitHub Enterprise

The instructions for GitHub Enterprise are similar, with a few extra steps to enable Enterprise and create the required default token.

In the case of GitHub Enterprise add the `defaultToken` created in the [prerequisite phase](#) to the `GitHub` section. The hostname should not include the protocol, ex: `github.exampleorg.com`.

**Option 1:** Create the Kubernetes Secret yourself.

```
kubectl create secret generic github-secret \
  --from-literal=clientId=<client-id> \
  --from-literal=clientSecret=<client-secret> \
  --from-literal=defaultToken=<default-token>
```

You must then provide the following to the `values.yaml` file:

```
github:
  enterprise: true
  hostname: "<github-enterprise-hostname>"
```

### Option 2: CircleCI creates the Kubernetes Secret.

Add `clientId`, `clientSecret` and `defaultToken` to the `values.yaml` file. You must also set `enterprise` to `true`, and provide the `hostname` for your enterprise GitHub. CircleCI will create the Kubernetes Secret automatically.

```
github:
  ...
  clientId: "<client-id>"
  clientSecret: "<client-secret>"
  enterprise: true
  hostname: "<github-enterprise-hostname>"
  defaultToken: "<token>"
```

## k. Object storage

Regardless of your storage provider, the bucket name you created during the [prerequisites phase](#) will need to be included.

```
object_storage:
  bucketName: "<bucket-name>"
```

### S3 compatible

Add an `s3` section as a child of `object_storage`. The `endpoint` in the case of AWS S3 is the [regional endpoint](#), it is of the form <https://s3.<region>.amazonaws.com>. Otherwise it is the API endpoint for your object storage server.

```
object_storage:
  ...
  s3:
    enabled: true
    endpoint: "<storage-server-or-s3-endpoint>"
```

Under `object_storage.s3`, you may provide the `accessKey` and `secretKey`, the `irsaRole`, or nothing. They were created during the prerequisites steps.

#### Option 1: Use IAM keys.

Add the following to the `object_storage.s3` section:

```
object_storage:
  ...
  s3:
    ...
    accessKey: "<access-key>"
    secretKey: "<secret-key>"
```

#### Option 2: Use IRSA.

A Kubernetes Secret will automatically be generated for you using your credentials.

Add the following to the `object_storage.s3` section:

```
object_storage:
  ...
  s3:
    ...
    region: "<role-region>"
    irsaRole: "<irsa-arn>"
```

#### Option 3: Create the Kubernetes Secret yourself

Instead of providing AWS `accessKey` and `secretKey` credentials in your `values.yaml` file, you may choose to create the Kubernetes Secret yourself.

```
kubectl create secret generic object-storage-secret \
  --from-literal=s3AccessKey=<access-key> \
  --from-literal=s3SecretKey=<secret-key>
```

CircleCI server will use the role provided to authenticate to S3.

## I. Installing behind a proxy

Depending on your security requirements, you might want to install CircleCI server behind a proxy. Installing behind a proxy gives you the power to monitor and control access between your installation and the broader Internet. For further information including limitations of installation behind a proxy, see the [Installing Server Behind a Proxy](#) guide.

The following fields need to be configured in your `values.yaml`:

- Toggle `proxy.enabled` to `"1"`
- Enter details for `proxy.http.host` and `proxy.https.host`, along with their associated ports. These values can be the same but they both need to be configured.
- For authentication you will need to configure `proxy.http.auth.enabled` and `proxy.https.auth.enabled` as `"1"`. You will also need to configure the respective username and password for both HTTP and HTTPS.
- configure the `no_proxy` hosts and subnets. This should include `localhost`, your GitHub Enterprise host (optional), the hostname of your CircleCI installation (see [Known Limitations](#) for an explanation), and the CIDRs of both `vm-service` and `Nomad`.

```
proxy:
  enabled: "1"
  http:
    host: "<proxy.example.internal>"
    port: "3128"
    auth:
      enabled: "1"
      username: "<proxy-user>"
      password: "<proxy-password>"
  https:
    host: "<proxy.example.internal>"
    port: "3128"
    auth:
      enabled: "1"
      username: "<proxy-user>"
      password: "<proxy-password>"
  no_proxy:
    - localhost
    - 127.0.0.1
    - "<github.example.internal>"
    - "<circleci.example.internal>"
    - "<nomad-subnet-cidr>"
    - "<vm-service-cidr>"
    - "<vpc-or-subnet-cidr>" # VPC or subnets to exclude from the proxy (optional)
```

## 4. Deploy

Once you have completed the fields detailed above, you can deploy CircleCI's core services:

```
USERNAME=<provided-username>
PASSWORD=<token>
namespace=<your-namespace>
helm registry login cciserver.azurecr.io/circleci-server -u $USERNAME -p $PASSWORD
helm install circleci-server oci://cciserver.azurecr.io/circleci-server -n $namespace --version
4.0.0 -f <path-to-values.yaml>
```

## 5. Create DNS entry

Create a DNS entry for your NGINX load balancer, for example, `circleci.your.domain.com` and `app.circleci.your.domain.com`. The DNS entry should align with the DNS names used when creating your TLS certificate and GitHub OAuth app during the prerequisites steps. All traffic will be routed through this DNS record.

You need the IP address, or, if using AWS, the DNS name of the NGINX load balancer. You can find this information with the following command:

```
kubectl get service circleci-proxy
```

For more information on adding a new DNS record, see the following documentation:

- [Managing Records \(GCP\)](#)
- [Creating records by using the Amazon Route 53 Console \(AWS\)](#)

## 6. Validation

You should now be able to navigate to your CircleCI server installation and log in to the application successfully.

Now we will move on to build services. It may take a while for all your services to be up. You can periodically check by running the following command (you are looking for the `frontend` pod to show a status of `running` and `ready` should show 1/1):

```
kubectl get pods -n <YOUR_CIRCLECI_NAMESPACE>
```



VM service and Nomad server pods are expected to fail at this stage. You will set up your execution environments in the next phase of the installation.

# Phase 3 - Execution environments

Before you begin with the CircleCI server v4.x execution environment installation phase, ensure you have run through [Phase 1 – Prerequisites](#) and [Phase 2 - Core Services Installation](#).



In the following sections, replace any sections indicated by `< >` with your details.

## 1. Nomad clients

Nomad is a workload orchestration tool that CircleCI uses to schedule (through Nomad server) and run (through Nomad clients) CircleCI jobs.

Nomad clients are installed outside of the Kubernetes cluster, while their control plane (Nomad Server) is installed within the cluster. Communication between your Nomad Clients and the Nomad control plane is secured with mTLS. The mTLS certificate, private key, and certificate authority will be output after you complete installation of the Nomad Clients.

### a. Create your cluster with Terraform

CircleCI curates Terraform modules to help install Nomad clients in your chosen cloud provider. You can browse the modules in our [public repository](#), including example Terraform config files for both AWS and GCP.

#### AWS cluster

You need some information about your cluster and server installation to populate the required variables for the Terraform module. A full example, as well as a full list of variables, can be found in the [example AWS Terraform configuration](#).

- **Server\_endpoint** - This is the domain name of the CircleCI application appended with port number `4647`.
- **Subnet ID (subnet)**, **VPC ID (vpc\_id)**, and **DNS server (dns\_server)** of your cluster. Run the following command to get the cluster VPC ID (vpclId), and subnets (subnetIds):

```
# Fetch VPC ID
aws eks describe-cluster --name=<cluster-name> --query "cluster.resourcesVpcConfig.vpcId"
--region=<region> --output text | xargs

# Fetch Subnet IDs
aws eks describe-cluster --name=<cluster-name> --query "cluster.resourcesVpcConfig.subnetIds"
--region=<region> --output text | xargs
```

This returns something similar to the following:

```
# VPC Id
vpc-02fdfff4ca

# Subnet Ids
subnet-08922063f12541f93 subnet-03b94b6fb1e5c2a1d subnet-0540dd7b2b2ddb57e subnet-
01833e1fa70aa4488
```

Then, using the VPCID you just found, run the following command to get the CIDR Block for your cluster. For AWS, the DNS Server is the third IP in your CIDR block (`CidrBlock`), for example your CIDR block might be `10.100.0.0/16`, so the third IP would be `10.100.0.2`.

```
aws ec2 describe-vpcs --filters Name=vpc-id,Values=<vpc-id> --query "Vpcs[0].CidrBlock" --region
=<region> --output text | xargs
```

This returns something like the following:

```
192.168.0.0/16
```

Once you have filled in the appropriate information, you can deploy your Nomad clients by running the following commands:

```
terraform init
```

```
terraform plan
```

```
terraform apply
```

After Terraform is done spinning up the Nomad client(s), it outputs the certificates and keys needed for configuring the Nomad control plane in CircleCI server. Copy them somewhere safe. The apply process usually only takes a minute.

## b. Nomad Autoscaler configuration

Nomad can automatically scale up or down your Nomad clients, provided your clients are managed by a cloud provider's autoscaling resource. With Nomad Autoscaler, you need to provide permission for the utility to manage your autoscaling resource and specify where it is located. CircleCI's Nomad terraform module can provision the permissions resources, or it can be done manually.

### AWS autoscaler IAM/role

Create an IAM user or role and policy for Nomad Autoscaler. You may take **one** of the following approaches:



- The CircleCI [nomad module](#) creates an IAM user and outputs the keys if you set variable `nomad_auto_scaler = true`. You may reference the example in the link for more details. If you have already created the clients, you can update the variable and run `terraform apply`. The created user's access and secret key will be available in Terraform's output.
- Create a Nomad Autoscaler IAM user manually with the [IAM policy below](#). Then, generate an access and secret key for this user.
- You may create a [Role for Service Accounts](#) for Nomad Autoscaler and attach the [IAM policy below](#):

When using access and secret keys, you have two options for configuration:

**Option 1:** CircleCI creates the Kubernetes Secret for you.

Add the access key and secret to your `values.yaml` file with any additional required configuration:

```
nomad:
  ...
  auto_scaler:
    aws:
      accessKey: "<access-key>"
      secretKey: "<secret-key>"
```

**Option 2:** Create the Kubernetes Secret yourself

Instead of storing the access key and secret in your `values.yaml` file, you may create the Kubernetes Secret yourself.



When using this method, an additional field is required for this secret, as outlined below.

First, add your access key, secret, and region to the following text, and encode it all with base64.

```
ADDITIONAL_CONFIG=`cat << EOF | base64
target "aws-asg" {
  driver = "aws-asg"
  config = {
    aws_region = "<aws-region>"
    aws_access_key_id = "<access-key>"
    aws_secret_access_key = "<secret-key>"
  }
}
EOF`
```

Then, using that additional base64 encoded config, create the Kubernetes Secret.

```
# With the base64-encoded additional config from above
kubectl create secret generic nomad-autoscaler-secret \
  --from-literal=secret.hcl=$ADDITIONAL_CONFIG
```

#### IAM policy for Nomad Autoscaler

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Sid": "VisualEditor0",
      "Effect": "Allow",
      "Action": [
        "autoscaling:CreateOrUpdateTags",
        "autoscaling:UpdateAutoScalingGroup",
        "autoscaling:TerminateInstanceInAutoScalingGroup"
      ],
      "Resource": "<<Your Autoscaling Group ARN>>"
    },
    {
      "Sid": "VisualEditor1",
      "Effect": "Allow",
      "Action": [
        "autoscaling:DescribeScalingActivities",
        "autoscaling:DescribeAutoScalingGroups"
      ],
      "Resource": "*"
    }
  ]
}
```

## 2. Nomad Servers

Now that you have successfully deployed your Nomad clients and have the permission resources, you can configure the Nomad Servers.

### a. Nomad gossip encryption key

Nomad requires a key to encrypt communications. This key must be exactly 32 bytes long. CircleCI will not be able to recover the values if lost. Depending on how you prefer to manage Kubernetes Secrets, there are two options:

**Option 1:** Create the Kubernetes Secret yourself.

```
kubectl create secret generic nomad-gossip-encryption-key \
--from-literal=gossip-key=<secret-key-32-chars>
```

Once the Kubernetes Secret exists, no change to `values.yaml` is required. The Kubernetes Secret will be referenced by default.

**Option 2:** CircleCI creates the Kubernetes Secret.

Add the value to `values.yaml`. CircleCI will create the Kubernetes Secret automatically.

```
nomad:
  server:
    gossip:
      encryption:
        key: "<secret-key-32-chars>"
```

## b. Nomad mTLS

The `CACertificate`, `certificate` and `privateKey` can be found in the output of the terraform module. They must be base64 encoded.

```
nomad:
  server:
    ...
    rpc:
      mTLS:
        enabled: true
        certificate: "<base64-encoded-certificate>"
        privateKey: "<base64-encoded-private-key>"
        CACertificate: "<base64-encoded-ca-certificate>"
```

## c. Nomad Autoscaler

If you have enabled Nomad Autoscaler, also include the following section under `nomad`:

### AWS

You created these values in the [Nomad Autoscaler Configuration section](#).

```

nomad:
  ...
  auto_scaler:
    enabled: true
    scaling:
      max: <max-node-limit>
      min: <min-node-limit>

  aws:
    enabled: true
    region: "<region>"
    autoScalingGroup: "<asg-name>"

    accessKey: "<access-key>"
    secretKey: "<secret-key>"
    # or
    irsaRole: "<role-arn>"

```

## d. Helm upgrade

Apply the changes made to your `values.yaml` file:

```

namespace=<your-namespace>
helm upgrade circleci-server oci://cciserver.azurecr.io/circleci-server -n $namespace --version
4.0.0 -f <path-to-values.yaml>

```

## 3. Nomad Clients Validation

CircleCI has created a project called [realitycheck](#) which allows you to test your server installation. We are going to follow the project so we can verify that the system is working as expected. As you continue through the next phase, sections of realitycheck will move from red (fail) to green (pass).

Before running realitycheck, check if the Nomad servers can communicate with the Nomad clients by executing the below command.

```

kubectl -n <namespace> exec -it $(kubectl -n <namespace> get pods -l app=nomad-server -o name | tail
-1) -- nomad node status

```

You should be able to see output like this:

ID	DC	Name	Class	Drain	Eligibility	Status
132ed55b	default	ip-192-168-44-29	linux-64bit	false	eligible	ready

To run realitycheck, you need to clone the repository. Depending on your GitHub setup, you can use one of the following commands:

#### GitHub Cloud

```
git clone https://github.com/circleci/realitycheck.git
```

#### GitHub Enterprise

```
git clone https://github.com/circleci/realitycheck.git
git remote set-url origin <YOUR_GH_REPO_URL>
git push
```

Once you have successfully cloned the repository, you can follow it from within your CircleCI server installation. You need to set the following variables. For full instructions please see the [repository readme](#).

Table 1. Environmental Variables

Name	Value
CIRCLE_HOSTNAME	<YOUR_CIRCLECI_INSTALLATION_URL>
CIRCLE_TOKEN	<YOUR_CIRCLECI_API_TOKEN>

Table 2. Contexts

Name	Environmental Variable Key	Environmental Variable Value
org-global	CONTEXT_END_TO_END_TEST_VAR	Leave blank
individual-local	MULTI_CONTEXT_END_TO_END_VAR	Leave blank

Once you have configured the environmental variables and contexts, rerun the realitycheck tests. You should see the features and resource jobs complete successfully. Your test results should look something like the following:

Pipeline	Status	Workflow	Branch / Commit	Start	Duration	Actions
realitycheck 3	Running	vm_jobs	master 368b13d	19s ago	1m 48s	
Jobs						
	remote_docker 60				1m 19s	
	machine_dlc 62				1m 29s	
	machine 63				1m 23s	
	docker_layer_caching 61				1m 19s	
	Success	resource_class_jobs	master 368b13d	29s ago	29s	
Jobs						
	xlarge 55				20s	
	small 58				2s	
	medium+ 57				20s	
	medium 59				3s	
	large 56				21s	
	Success	feature_jobs	master 368b13d	2m ago	31s	
Jobs						
	write_workspace 51				6s	
	save_and_restore_cache 49				27s	
	read_workspace 54				2s	
	multi-contexts 52				6s	
	contexts 53				5s	
	artifacts_test_results 50				7s	

### 3. VM service

VM service configures virtual machine and remote docker jobs. You can configure a number of options for VM service, such as scaling rules. VM service is unique to AWS and GCP installations because it relies on specific features of these cloud providers.

#### AWS

##### Set up security group

###### 1. Get the information needed to create security groups

The following command returns your VPC ID (`vpcId`) and CIDR Block (`serviceIpv4Cidr`) which you need throughout this section:

```
# Fetch VPC Id
aws eks describe-cluster --name=<cluster-name> --query "cluster.resourcesVpcConfig.vpcId"
--region=<region> --output text | xargs

# Fetch CIDR Block
aws eks describe-cluster --name=<cluster-name> --query
"cluster.kubernetesNetworkConfig.serviceIpv4Cidr" --region=<region> --output text | xargs
```

###### 2. Create a security group

Run the following commands to create a security group for VM service:

```
aws ec2 create-security-group --vpc-id "<VPC_ID>" --description "CircleCI VM Service security group" --group-name "circleci-vm-service-sg"
```

This outputs a GroupID to be used in the next steps:

```
{
  "GroupId": "<VM_SECURITY_GROUP_ID>"
}
```

### 3. Apply security group Nomad

Use the security group you just created, and your CIDR block values, to apply the security group. This allows VM service to communicate with created EC2 instances on port 22.

```
aws ec2 authorize-security-group-ingress --group-id "<VM_SECURITY_GROUP_ID>" --protocol tcp --port 22 --cidr "<SERVICE_IPV4_CIDR>"
```

For each [subnet](#) used by the Nomad clients, find the subnet cidr block and add two rules with the following commands.

```
# find CIDR block
aws ec2 describe-subnets --subnet-ids=<nomad-subnet-id> --query "Subnets[*].[SubnetId, CidrBlock]" --region=<region> | xargs
```

```
# add a security group allowing docker access from nomad clients, to VM instances
aws ec2 authorize-security-group-ingress --group-id "<VM_SECURITY_GROUP_ID>" --protocol tcp --port 2376 --cidr "<SUBNET_IPV4_CIDR>"
```

```
# add a security group allowing SSH access from nomad clients, to VM instances
aws ec2 authorize-security-group-ingress --group-id "<VM_SECURITY_GROUP_ID>" --protocol tcp --port 22 --cidr "<SUBNET_IPV4_CIDR>"
```

### 4. Apply the security group for SSH (If using public IPs for machines)

If using public IPs for VM service instances, run the following command to apply the security group rules so users can SSH into their jobs:

```
aws ec2 authorize-security-group-ingress --group-id "<VM_SECURITY_GROUP_ID>" --protocol tcp --port 54782 --cidr "0.0.0.0/0"
```

## Set up authentication

There are two ways to authenticate CircleCI with your cloud provider: IAM Roles for Service Accounts (IRSA), and IAM access keys. IRSA is the recommended method.

The following is a summary of [AWS's Documentation on IRSA](#) that is sufficient for configuring VM Service in CircleCI.

### 1. Create indentity provider

Create an IAM OIDC identity provider for your EKS Cluster:

```
eksctl utils associate-iam-oidc-provider --cluster <CLUSTER_NAME> --approve
```

### 2. Get ARN

Get the OIDC provider ARN with the following command, you will need it in later steps:

```
aws iam list-open-id-connect-providers | grep $(aws eks describe-cluster --name <CLUSTER_NAME> --query "cluster.identity.oidc.issuer" --output text | awk -F '/' '{print $NF}')
```

### 3. Get URL

Get your OIDC provider URL, you will need it in later steps

```
aws eks describe-cluster --name <CLUSTER_NAME> --query "cluster.identity.oidc.issuer" --output text | sed -e "s/^https:\/\/\/"
```

### 4. Create role

Create the role using the command and trust policy template below, you will need the Role ARN and name in later steps:

```
aws iam create-role --role-name circleci-vm --assume-role-policy-document file:///<TRUST_POLICY_FILE>
```



```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "Federated": "<OIDC_PROVIDER_ARN>"
      },
      "Action": "sts:AssumeRoleWithWebIdentity",
      "Condition": {
        "StringEquals": {
          "<OIDC_PROVIDER_URL>:sub": "system:serviceaccount:<K8S_NAMESPACE>:vm-service"
        }
      }
    }
  ]
}
```

## 5. Create policy

Create the policy using the command and template below. Fill in the security group ID and the VPC ID:

```
aws iam create-policy --policy-name circleci-vm --policy-document file:///<POLICY_FILE>
```

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": "ec2:RunInstances",
      "Effect": "Allow",
      "Resource": [
        "arn:aws:ec2::*:image/*",
        "arn:aws:ec2::*:snapshot/*",
        "arn:aws:ec2::*:key-pair/*",
        "arn:aws:ec2::*:launch-template/*",
        "arn:aws:ec2::*:network-interface/*",
        "arn:aws:ec2::*:placement-group/*",
        "arn:aws:ec2::*:volume/*",
        "arn:aws:ec2::*:subnet/*",
        "arn:aws:ec2::*:security-group/<SECURITY_GROUP_ID>"
      ]
    }
  ],
}
```

```

{
  "Action": "ec2:RunInstances",
  "Effect": "Allow",
  "Resource": "arn:aws:ec2:*:*:instance/*",
  "Condition": {
    "StringEquals": {
      "aws:RequestTag/ManagedBy": "circleci-vm-service"
    }
  }
},
{
  "Action": [
    "ec2:CreateVolume"
  ],
  "Effect": "Allow",
  "Resource": [
    "arn:aws:ec2:*:*:volume/*"
  ],
  "Condition": {
    "StringEquals": {
      "aws:RequestTag/ManagedBy": "circleci-vm-service"
    }
  }
},
{
  "Action": [
    "ec2:Describe*"
  ],
  "Effect": "Allow",
  "Resource": "*"
},
{
  "Effect": "Allow",
  "Action": [
    "ec2:CreateTags"
  ],
  "Resource": "arn:aws:ec2:*:*:*/*/*",
  "Condition": {
    "StringEquals": {
      "ec2:CreateAction" : "CreateVolume"
    }
  }
}

```

```

    }
  },
  {
    "Effect": "Allow",
    "Action": [
      "ec2:CreateTags"
    ],
    "Resource": "arn:aws:ec2:*:*:*/*",
    "Condition": {
      "StringEquals": {
        "ec2:CreateAction" : "RunInstances"
      }
    }
  },
  {
    "Action": [
      "ec2:CreateTags",
      "ec2:StartInstances",
      "ec2:StopInstances",
      "ec2:TerminateInstances",
      "ec2:AttachVolume",
      "ec2:DetachVolume",
      "ec2>DeleteVolume"
    ],
    "Effect": "Allow",
    "Resource": "arn:aws:ec2:*:*:*/*",
    "Condition": {
      "StringEquals": {
        "ec2:ResourceTag/ManagedBy": "circleci-vm-service"
      }
    }
  },
  {
    "Action": [
      "ec2:RunInstances",
      "ec2:StartInstances",
      "ec2:StopInstances",
      "ec2:TerminateInstances"
    ],
    "Effect": "Allow",
    "Resource": "arn:aws:ec2:*:*:subnet/*",

```

```

    "Condition": {
      "StringEquals": {
        "ec2:Vpc": "<VPC_ID>"
      }
    }
  }
]
}

```

## 6. Attach policy

Attach the policy to the role:

```
aws iam attach-role-policy --role-name <VM_ROLE_NAME> --policy-arn=<VM_POLICY_ARN>
```

## 7. Configure VM-Service

Configure VM service by adding the following to `values.yaml`:

```

vm_service:
  providers:
    ec2:
      enabled: true
      region: "<REGION>"
      assignPublicIP: true
      irsaRole: "<IRSA_ROLE_ARN>"
      subnets:
        - "<SUBNET_ID>"
      securityGroupId: "<SECURITY_GROUP_ID>"

```

## 1. Create user

Create a new user with programmatic access:

```
aws iam create-user --user-name circleci-vm-service
```

Optionally, vm-service does support the use of a [service account role](#) in place of AWS keys. If you would prefer to use a role, follow these [instructions](#) using the policy in step 6 below. Once done, you may skip to step 9, enabling vm-service.

## 2. Create policy

Create a `policy.json` file with the following content. You should fill in the ID of the VM Service security group created in step 2 (`VMServiceSecurityGroupId`) and VPC ID (`vpcID`) below.

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Action": "ec2:RunInstances",
      "Effect": "Allow",
      "Resource": [
        "arn:aws:ec2:*::image/*",
        "arn:aws:ec2:*::snapshot/*",
        "arn:aws:ec2:*::key-pair/*",
        "arn:aws:ec2:*::launch-template/*",
        "arn:aws:ec2:*::network-interface/*",
        "arn:aws:ec2:*::placement-group/*",
        "arn:aws:ec2:*::volume/*",
        "arn:aws:ec2:*::subnet/*",
        "arn:aws:ec2:*::security-group/<YOUR_VMSecurityGroupId>"
      ]
    },
    {
      "Action": "ec2:RunInstances",
      "Effect": "Allow",
      "Resource": "arn:aws:ec2:*::instance/*",
      "Condition": {
        "StringEquals": {
          "aws:RequestTag/ManagedBy": "circleci-vm-service"
        }
      }
    },
    {
      "Action": [
        "ec2:CreateVolume"
      ],
      "Effect": "Allow",
      "Resource": [
        "arn:aws:ec2:*::volume/*"
      ],
      "Condition": {
        "StringEquals": {
          "aws:RequestTag/ManagedBy": "circleci-vm-service"
        }
      }
    }
  ]
}
```

```

    }
  },
  {
    "Action": [
      "ec2:Describe*"
    ],
    "Effect": "Allow",
    "Resource": "*"
  },
  {
    "Effect": "Allow",
    "Action": [
      "ec2:CreateTags"
    ],
    "Resource": "arn:aws:ec2:*:*:*/*/*",
    "Condition": {
      "StringEquals": {
        "ec2:CreateAction" : "CreateVolume"
      }
    }
  },
  {
    "Effect": "Allow",
    "Action": [
      "ec2:CreateTags"
    ],
    "Resource": "arn:aws:ec2:*:*:*/*/*",
    "Condition": {
      "StringEquals": {
        "ec2:CreateAction" : "RunInstances"
      }
    }
  },
  {
    "Action": [
      "ec2:CreateTags",
      "ec2:StartInstances",
      "ec2:StopInstances",
      "ec2:TerminateInstances",
      "ec2:AttachVolume",
      "ec2:DetachVolume",

```

```

        "ec2:DeleteVolume"
    ],
    "Effect": "Allow",
    "Resource": "arn:aws:ec2:*:*:*/*",
    "Condition": {
        "StringEquals": {
            "ec2:ResourceTag/ManagedBy": "circleci-vm-service"
        }
    }
},
{
    "Action": [
        "ec2:RunInstances",
        "ec2:StartInstances",
        "ec2:StopInstances",
        "ec2:TerminateInstances"
    ],
    "Effect": "Allow",
    "Resource": "arn:aws:ec2:*:*:subnet/*",
    "Condition": {
        "StringEquals": {
            "ec2:Vpc": "<vpc-id>"
        }
    }
}
]
}

```

### 3. Attach policy to user

Once you have created the policy.json file, attach it to an IAM policy and created user:

```

aws iam put-user-policy --user-name circleci-vm-service --policy-name circleci-vm-service
--policy-document file://policy.json

```

### 4. Create an access key and secret for the user

If you have not already created them, you will need an access key and secret for the `circleci-vm-service` user. You can create those by running the following command:

```

aws iam create-access-key --user-name circleci-vm-service

```

## 5. Configure server (there are two options)

- **Option 1 - Add the keys to `values.yaml`** Add the VM Service configuration to `values.yaml`. Details of the available configuration options can be found in the [Managing Virtual Machines with VM Service](#) guide.
- **Option 2 - Create the Kubernetes Secret yourself** Instead of providing the access key and secret in your `values.yaml` file, you may create the Kubernetes Secret yourself.

```
kubectl create secret generic vm-service-secret \
  --from-literal=ec2AccessKey=<access-key> \
  --from-literal=ec2SecretKey=<secret-key>
```

## VM Service Validation

Apply the changes made to your `values.yaml` file.

```
namespace=<your-namespace>
helm upgrade circleci-server oci://cciserver.azurecr.io/circleci-server -n $namespace --version
4.0.0 -f <path-to-values.yaml>
```

Once you have configured and deployed CircleCI server, you should validate that VM Service is operational. You can rerun the realitychecker project within your CircleCI installation and you should see the VM Service Jobs complete. At this point, all tests should pass.

## 4. Runner

### Overview

CircleCI runner does not require any additional server configuration. Server ships ready to work with runner. However, you need to create a runner and configure the runner agent to be aware of your server installation. For complete instructions for setting up runner, see the [runner documentation](#).



Runner requires a namespace per organization. Server can have many organizations. If your company has multiple organizations within your CircleCI installation, you need to set up a runner namespace for each organization within your server installation.



# Phase 4 - Post installation

Before you begin with the CircleCI server v4.x post installation phase, ensure you have run through [Phase 1 - Prerequisites](#), [Phase 2 - Core Services Installation](#) and [Phase 3 - Execution Environments Installation](#).



In the following sections, replace any sections indicated by `< >` with your details.

## Backup and restore

Backups of CircleCI server can be created through [Velero](#). You installed Velero in your cluster during the prerequisites installation phase.

### Set up backup and restore on AWS

These instructions were sourced from the [Velero docs](#).

#### 1. Create an AWS S3 bucket

```
BUCKET=<YOUR_BUCKET>
REGION=<YOUR_REGION>
aws s3api create-bucket \
  --bucket $BUCKET \
  --region $REGION \
  --create-bucket-configuration LocationConstraint=$REGION
```



`us-east-1` does not support a [LocationConstraint](#). If your region is `us-east-1`, omit the bucket configuration.

#### 2. Set up permissions for Velero

Create an IAM user:

```
aws iam create-user --user-name velero
```

Attach policies to give user `velero` the necessary permissions:

```
cat > velero-policy.json <<EOF
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
```

```

        "ec2:DescribeVolumes",
        "ec2:DescribeSnapshots",
        "ec2:CreateTags",
        "ec2:CreateVolume",
        "ec2:CreateSnapshot",
        "ec2>DeleteSnapshot"
    ],
    "Resource": "*"
},
{
    "Effect": "Allow",
    "Action": [
        "s3:GetObject",
        "s3:DeleteObject",
        "s3:PutObject",
        "s3:AbortMultipartUpload",
        "s3:ListMultipartUploadParts"
    ],
    "Resource": [
        "arn:aws:s3:::${BUCKET}/*"
    ]
},
{
    "Effect": "Allow",
    "Action": [
        "s3:ListBucket"
    ],
    "Resource": [
        "arn:aws:s3:::${BUCKET}"
    ]
}
]
}
EOF

```

```

aws iam put-user-policy \
  --user-name velero \
  --policy-name velero \
  --policy-document file://velero-policy.json

```

Create an access key for user `velero`:

```
aws iam create-access-key --user-name velero
```

The result should look like this:

```
{
  "AccessKey": {
    "UserName": "velero",
    "Status": "Active",
    "CreateDate": "2017-07-31T22:24:41.576Z",
    "SecretAccessKey": <AWS_SECRET_ACCESS_KEY>,
    "AccessKeyId": <AWS_ACCESS_KEY_ID>
  }
}
```

Create a Velero-specific credentials file (for example: `./credentials-velero`) in your local directory, with the following contents:

```
[default]
aws_access_key_id=<AWS_ACCESS_KEY_ID>
aws_secret_access_key=<AWS_SECRET_ACCESS_KEY>
```



the `AWS_ACCESS_KEY_ID` and `AWS_SECRET_ACCESS_KEY` placeholders are values returned from the `create-access-key` request in the previous step.

### 3. Install and start Velero

Run the following `velero install` command. This creates a namespace called `velero` and installs all the necessary resources to run Velero. Make sure that you pass the correct file name containing the AWS credentials that you created in step two.



Backups require [restic](#) to operate. When installing Velero, ensure that you have the `--use-restic` flag set, as shown below:

```

velero install \
  --provider aws \
  --plugins velero/velero-plugin-for-aws:v1.2.0 \
  --bucket $BUCKET \
  --backup-location-config region=$REGION \
  --snapshot-location-config region=$REGION \
  --secret-file ./credentials-velero \
  --use-restic \
  --wait

```

#### 4. Verify Velero

Once Velero is installed on your cluster, check the new `velero` namespace. You should have a Velero deployment and a restic daemonset, for example:

```

$ kubectl get pods --namespace velero

```

NAME	READY	STATUS	RESTARTS	AGE
restic-5v1ww	1/1	Running	0	2m
restic-94ptv	1/1	Running	0	2m
restic-ch6m9	1/1	Running	0	2m
restic-mknws	1/1	Running	0	2m
velero-68788b675c-dm2s7	1/1	Running	0	2m



As restic is a daemonset, there should be one pod for each node in your Kubernetes cluster.

## Set up backup and restore with S3-compatible storage

The following steps assume you are using S3-compatible object storage, but not necessarily AWS S3, for your backups.

These instructions were sourced from the [Velero docs](#).

#### 1. Configure `mc` client

To start, [configure](#) `mc` to connect to your storage provider:

```

# Alias can be any name as long as you use the same value in subsequent commands
export ALIAS=my-provider
mc alias set $ALIAS <YOUR_MINIO_ENDPOINT> <YOUR_MINIO_ACCESS_KEY_ID>
<YOUR_MINIO_SECRET_ACCESS_KEY>

```

You can verify your client is correctly configured by running `mc ls my-provider` and you should see the buckets in your provider enumerated in the output.

## 2. Create a bucket

Create a bucket for your backups. It is important that a new bucket is used, as Velero cannot use a preexisting bucket that contains other content.

```
mc mb ${ALIAS}/<YOUR_BUCKET>
```

## 3. Create a user and policy

Create a user and policy for Velero to access your bucket.



In the following snippet `<YOUR_MINIO_ACCESS_KEY_ID>` and `<YOUR_MINIO_SECRET_ACCESS_KEY>` refer to the credentials used by Velero to access MinIO.

```
# Create user
mc admin user add $ALIAS <YOUR_MINIO_ACCESS_KEY_ID> <YOUR_MINIO_SECRET_ACCESS_KEY>

# Create policy
cat > velero-policy.json << EOF
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": [
        "s3:*"
      ],
      "Resource": [
        "arn:aws:s3:::<YOUR_BUCKET>",
        "arn:aws:s3:::<YOUR_BUCKET>/*"
      ]
    }
  ]
}
EOF

mc admin policy add $ALIAS velero-policy velero-policy.json

# Bind user to policy
mc admin policy set $ALIAS velero-policy user=<YOUR_VELERO_ACCESS_KEY_ID>
```

Finally, you add your new user's credentials to a file (`./credentials-velero` in this example) with the

following contents:

```
[default]
aws_access_key_id=<YOUR_VELERO_ACCESS_KEY_ID>
aws_secret_access_key=<YOUR_VELERO_SECRET_ACCESS_KEY>
```

#### 4. Install and start Velero

Run the following `velero install` command. This creates a namespace called `velero` and installs all the necessary resources to run Velero.



Backups require [restic](#) to operate. When installing Velero, ensure that you have the `--use-restic` flag set, as shown below:

```
velero install --provider aws \
  --plugins velero/velero-plugin-for-aws:v1.2.0 \
  --bucket <YOUR_BUCKET> \
  --secret-file ./credentials-velero \
  --use-volume-snapshots=false \
  --use-restic \
  --backup-location-config region=minio,s3ForcePathStyle="true",s3Url=<YOUR_ENDPOINT> \
  --wait
```

#### 5. Verify Velero

Once Velero is installed on your cluster, check the new `velero` namespace. You should have a Velero deployment and a restic daemonset, for example:

```
$ kubectl get pods --namespace velero
```

NAME	READY	STATUS	RESTARTS	AGE
restic-5vlww	1/1	Running	0	2m
restic-94ptv	1/1	Running	0	2m
restic-ch6m9	1/1	Running	0	2m
restic-mknws	1/1	Running	0	2m
velero-68788b675c-dm2s7	1/1	Running	0	2m



As restic is a daemonset, there should be one pod for each node in your Kubernetes cluster.

### Take a backup

Now that Velero is installed on your cluster, you are ready to create your first backup. If you encounter problems, please refer to the [troubleshooting](#) section.

- To create the backup, run:

```
K8S_NS=$(helm list -o yaml | yq '.[].namespace')
CHART=$(helm list -o yaml | yq '.[].chart' )
REV=$(helm list -o yaml | yq '.[].revision')
RANDOM_STR=$(cat /dev/urandom | env LC_ALL=C tr -dc 'a-z0-9' | head -c 8)

velero backup create "${K8S_NS}-${RANDOM_STR}" --include-namespaces "${K8S_NS}" --labels "chart--
rev=${CHART}--${REV}"
```

- To restore from a backup, run:

```
# List all existing backups
velero backup get --show-labels

# Restore the specific backup
velero restore create --include-namespaces <circleci-namespace> --from-backup <backup-name>
```

See the [Velero](#) documentation or more details.

## Email Notifications

Add email notification support by adding the following to `values.yaml`:

```
smtp:
  host: <hostname-of-submission-server>
  user: <username-for-submission-server>
  password: <password-for-submission-server>
  port: <mail-port>
```

## Managing orbs

Server installations include their own local orb registry. This registry is private to the server installation. All orbs referenced in project configs reference the orbs in the *server* orb registry. You are responsible for maintaining orbs. This includes:

- Copying orbs from the public registry.
- Updating orbs that may have been copied previously.
- Registering your company's private orbs, if you have any.

For more information and steps to complete these tasks, see the [Orbs on server guide](#).

# Hardening your cluster

This section provides supplemental information on hardening your Kubernetes cluster.

## Network topology

A server installation basically runs three different type of compute instances: The Kubernetes nodes, Nomad clients, and external VMs.

It is highly recommended that you deploy these into separate subnets with distinct CIDR blocks. This will make it easier for you to control traffic between the different components of the system and isolate them from each other.

Best practice is to make as many of the resources as private as possible. If your users will access your CircleCI server installation via VPN, there is no need to assign any public IP addresses at all, as long as you have a working NAT gateway setup. Otherwise, you will need at least one public subnet for the `circleci-proxy` load balancer.

However, in this case, it is also recommended to place Nomad clients and VMs in a public subnet to enable your users to SSH into jobs and scope access via networking rules.

Currently, custom subnetting is not supported for GCP. Custom subnetting support will be available in a future update/release.



An NGINX reverse proxy is placed infront of `Kong` and exposed as a Kubernetes service named `circleci-proxy`. NGINX is responsible routing the traffic to the following services: `kong`, `vm-service`, `output-processor` and `nomad`.



When using Amazon Certificate Manager (ACM), the name of nginx's service will be `circleci-proxy-acm` instead of `circleci-proxy`. If you have switched from some other method of handling your TLS certificates to using ACM, this change will recreate the loadbalancer and you will have to reroute your associated DNS records for your `<domain>` and `app.<domain>`.

## Network traffic

This section explains the minimum requirements for a server installation to work. Depending on your workloads, you might need to add additional rules to egress for Nomad clients and VMs. As nomenclature between cloud providers differs, you will probably need to implement these rules using firewall rules and/or security groups.

Where you see "external," this usually means all external IPv4 addresses. Depending on your particular setup, you might be able to be more specific (for example, if you are using a proxy for all external traffic).

The rules explained here are assumed to be stateful and for TCP connections only, unless stated otherwise. If you are working with stateless rules, you need to create matching ingress or egress rules for the ones listed here.



## Reverse proxy status

You may wish to check the status of the services routing traffic in your CircleCI server installation and alert if there are any issues. Since we use both Nginx and Kong in CircleCI server, we expose the status pages of both via port 80.

Service	Endpoint
nginx	/nginx_status
kong	/kong_status

## Kubernetes load balancers

Depending on your setup, your load balancers might be transparent (that is, they are not treated as a distinct layer in your networking topology). In this case, you can apply the rules from this section directly to the underlying destination or source of the network traffic. Refer to the documentation of your cloud provider to make sure you understand how to correctly apply networking security rules, given the type of load balancing you are using with your installation.

### Ingress

If the traffic rules for your load balancers have not been created automatically, here are their respective ports:

Name	Port	Source	Purpose
circleci-proxy/-acm	80	External	User Interface & Frontend API
circleci-proxy/-acm	443	External	User Interface & Frontend API
circleci-proxy/-acm	3000	Nomad clients	Communication with Nomad clients
circleci-proxy/-acm	4647	Nomad clients	Communication with Nomad clients
circleci-proxy/-acm	8585	Nomad clients	Communication with Nomad clients

### Egress

The only type of egress needed is TCP traffic to the Kubernetes nodes on the Kubernetes load balancer ports (30000-32767). This is not needed if your load balancers are transparent.

## Common rules for compute instances

These rules apply to all compute instances, but not to the load balancers.

## Ingress

If you want to access your instances using SSH, you will need to open port 22 for TCP connections for the instances in question. It is recommended to scope the rule as closely as possible to allowed source IPs and/or only add such a rule when needed.

## Egress

You most likely want all of your instances to access internet resources. This requires you to allow egress for UDP and TCP on port 53 to the DNS server within your VPC, as well as TCP ports 80 and 443 for HTTP and HTTPS traffic, respectively. Instances building jobs (that is, the Nomad clients and external VMs) also will likely need to pull code from your VCS using SSH (TCP port 22). SSH is also used to communicate with external VMs, so it should be allowed for all instances with the destination of the VM subnet and your VCS, at the very least.

## Kubernetes nodes

### Intra-node traffic

By default, the traffic within your Kubernetes cluster is regulated by networking policies. For most purposes, this should be sufficient to regulate the traffic between pods and there is no additional requirement to reduce traffic between Kubernetes nodes any further (it is fine to allow all traffic between Kubernetes nodes).

To make use of networking policies within your cluster, you may need to take additional steps, depending on your cloud provider and setup. Here are some resources to get you started:

- [Kubernetes Network Policy Overview](#)
- [Creating a Cluster Network Policy on Google Cloud](#)
- [Installing Calico on Amazon EKS](#)

## Ingress

If you are using a managed service, you can check the rules created for the traffic coming from the load balancers and the allowed port range. The standard port range for Kubernetes load balancers (30000-32767) should be all that is needed here for ingress. If you are using transparent load balancers, you need to apply the ingress rules listed for load balancers above.

## Egress

Port	Destination	Purpose
2376	VMs	Communication with VMs
4647	Nomad clients	Communication with the Nomad clients
all traffic	other nodes	Allow intra-cluster traffic

## Nomad Clients

Nomad clients do not need to communicate with each other. You can block traffic between Nomad client instances completely.

### Ingress

Port	Source	Purpose
4647	K8s nodes	Communication with Nomad server
64535-65535	External	Rerun jobs with SSH functionality

### Egress

Port	Destination	Purpose
2376	VMs	Communication with VMs
3000	VM Service load balancers	Internal communication
4647	Nomad Load Balancer	Internal communication
8585	Output Processor Load Balancer	Internal communication

## External VMs

Similar to Nomad clients, there is no need for external VMs to communicate with each other.

### Ingress

Port	Source	Purpose
22	Kubernetes nodes	Internal communication
22	Nomad clients	Internal communication
2376	Kubernetes nodes	Internal communication
2376	Nomad clients	Internal communication
54782	External	Rerun jobs with SSH functionality

### Egress

You will only need the egress rules for internet access and SSH for your VCS.

# Installing server behind a proxy

Depending on your security requirements, you might want to install CircleCI server behind a proxy. Installing behind a proxy gives you the power to monitor and control access between your installation and the broader Internet.

Configuring a proxy happens during [Phase 2 - Core services](#).

## Known limitations

- The CircleCI hostname must be added to the no-proxy list because the following services: `output processor` and `vm-service`. The application and build-agent share a no-proxy list and are assumed to be behind the same firewall and therefore cannot have a proxy between them.
- Some additional configuration is required to import orbs when installed behind a proxy. See [Orbs on Server](#) docs for more information.
- The JVM only accepts proxies that run over HTTP, not HTTPS, and therefore proxy URLs must be of the form `http://user:password@host:port` rather than `https://user:password@host:port`.
- If your GitHub instance is running outside of the proxied environment (either GitHub.com or GitHub Enterprise), you must ensure that SSH traffic from CircleCI (inside the Kubernetes cluster) and from our Nomad node can reach your instance. Please note the default `checkout` step in a CircleCI job will fail to clone code and our `ssh-keyscan` of GitHub Enterprise will not work. While you may configure an SSH proxy, `ssh-keyscan` can NOT be proxied and instead will require you provide `github.fingerprint` when using GHE.
- If you install server behind a proxy, you may need to provide a custom image for VM service. Visit the [CircleCI Linux Image Builder repo](#) for further information.
- If object storage is outside the proxy, no job features that use object storage will work. This includes:
  - Artifacts
  - Test results
  - Cache save and restore
  - Workspaces

Users can get around this restriction by setting environment variables on their jobs. For example:

```
jobs:
  my-job:
    docker:
      - image: cimg/node:17.2.0
    environment:
      HTTP_PROXY: http://proxy.example.com:3128
      HTTPS_PROXY: http://proxy.example.com:3128
      NO_PROXY: whatever.internal,10.0.1.2
```



It is crucial that these environment variables are set in this specific location because it is the only location that propagates them to the correct service.

# Migrate from server v3.x to v4.x

Migrating from server v3.x to v4.x is an *in-place* migration. You will generate a helm-value file, which will be used in a helm command to upgrade the CircleCI server from v3.x to v4.x.

We recommend using a staging environment before completing this process in a production environment. This will allow you to gain a better understanding of the migration process, and give you an idea of how long the migration will take to complete.

## Prerequisites

- Your current CircleCI server installation is v3.x.
- You have taken a backup of the v3.x instance, following the instructions at the [Backup and Restore](#) guide.
  - If you are using external datastores, they need to be backed up separately.
  - If you are using internal datastores they are backed up using the backup process described in the link above.
- The migration script must be run from a machine with following tools installed:
  - [kubect](#)l configured for the server 3.x instance
  - [yq](#)
  - [helm](#)
  - [helm-diff](#)
  - [Azure CLI](#)

## Migration

To perform an in-place upgrade of CircleCI server from v3.x to v4.x, you will execute a migration script, which will perform the following tasks:

- Export the KOTS config values into a YAML file.
- Modify the exported YAML file to be used with `helm` to migrate to server v4.x.
- Annotate all the kubernetes resources for `helm`.
- Execute postgres db `domain` migration.
- Execute cleanup steps which will delete all the KOTS related resources.
- Display an output message with the next steps.



In CircleCI server v4.x, we have changed the way we store workflow data. Previously we had stored your workflow data in Postgres. We have moved this data into your storage bucket (GCS/S3/MinIO). After migrating, there will be a period where your workflow data will be unavailable. This is due to your existing workflow data being transferred to your storage bucket. The time taken depends on how much existing workflow data you have, however, you may scale the `workflow-conductor` pods that are responsible for the migration to speed up the process. You may scale as much as your cluster will allow. In later steps we will show you how to scale your pods.

The `kots-exporter` script can also perform the following functions if required:

- Rerun the annotation step only (`-f annotate`)
- Run the kots-annotation cleanup job (`-f kots_cleanup`)
- Display the output message (`-f message`)
- Rerun the PostgreSQL database migration (`-f flyway`)

## 1. Create docker secret for CircleCI image registry

The migration script requires access to the CircleCI image registry to run the migration function. The image registry is also used in the server v4.x installation. Create the `docker-registry` secret in the same Kubernetes namespace in-which CircleCI server is installed. Image registry credentials will be provided by your CircleCI point of contact.

In a terminal run the following, substituting all values indicated by `< >`:

```
kubectl -n <namespace> create secret docker-registry regcred \
  --docker-server=https://cciserver.azurecr.io \
  --docker-username=<image-registry-username> \
  --docker-password=<image-registry-token> \
  --docker-email=<notification-email-id>
```

## 2. Clone the repository and run the kots-exporter script

The instructions below guide you through cloning the repository containing the KOTS exporter script. This script generates the `helm-value` file for server 4.x.

1. Run `git clone https://github.com/CircleCI-Public/server-scripts`.
2. Change into the `kots-exporter` directory: `cd server-scripts/kots-exporter`.
3. Run the migration script: `./kots-exporter.sh`.
4. You will be prompted for the following information:
  - Release name of your server 3.x installation - `circleci-server`
  - Kubernetes namespace of your server 3.x installation
  - License String for Server 4.x, which will be provided by your CircleCI point of contact.

5. After the script has completed, you will get a message detailing the next steps to perform, for example:

```
##### Upgrade Commands #####
Before upgrading to 4.0, follow below steps:

Your helm values file with your installation config is here:
- <path>/kots-exporter/output/helm-values.yaml
-----

## Progres Chart Upgrade Preparation
Upgrading to CircleCI server v4.x includes upgrading the Postgres chart.
Before upgrading, we need to prepare your Postgres instance.
This is only needed if your Postgres instance is not externalized.

# Collect your postgres user's password and the PVC attached to your postgres instance
export POSTGRESQL_PASSWORD=$(kubectl get secret --namespace <namespace> postgresql -o jsonpath
="{.data.postgresql-password}" | base64 --decode)
export POSTGRESQL_PVC=$(kubectl get pvc --namespace <namespace> -l app.kubernetes.io/instance
=circleci-server,role=primary -o jsonpath="{.items[0].metadata.name}")
# remove the postgres statefulset without terminating your postgres instance
kubectl delete statefulsets.apps postgresql --namespace <namespace> --cascade=orphan
# remove the existing secret containing your postgres password. This will get recreated during
upgrade.
kubectl delete secret postgresql --namespace <namespace>
-----

## Helm Diff
The Helm Diff tool is used to verify that the changes between your current install and the upgrade
are expected

# diff command
helm diff upgrade <release-name> -n <namespace> -f <path>/kots-exporter/output/helm-values.yaml
--show-secrets --context 5 oci://cciserver.azurecr.io/circleci-server --version 4.0.0
-----

## Helm Upgrade CircleCI Server
helm upgrade <release-name> -n <namespace> -f <path>/kots-exporter/output/helm-values.yaml
oci://cciserver.azurecr.io/circleci-server --version <version-to-upgrade-to> --force

NOTE: After server 3.x to 4.x migration, You must rerun the Nomad terraform with modified value of
'server_endpoint' variable
It should be - <domain-name>:4647
```



If your Postgres instance is not externalized, when upgrading to CircleCI Server v4.0, the Postgres chart is also being upgraded. Please note the command needed before running helm upgrade above.

### 3. Validate your helm-value file

Once the migration script has completed, a `helm-values.yaml` file will be generated with your existing CircleCI Server 3.x config. This file holds the configuration data you had previously entered in KOTS. Going forward, you will use this file to update/configure your CircleCI Server installation as is standard helm practice.



## 4. Generate helm-diff output

Next, execute the helm-diff command and review the output.

```
helm diff upgrade <release-name> -n <namespace> -f <path>/kots-exporter/output/helm-values.yaml  
--show-secrets --contexts 5 oci://cciserver.azurecr.io/circleci-server --version 4.0.0
```

Review the output generated from the helm-diff command using the following to help:

- line highlighted Yellow: Kubernetes resources status, for example, changed, added
- line highlighted Red: Deletion, for example, image
- line displayed in Green: Addition, for example, imagePullSecret

Below are the changes you should expect to see in helm-diff output:

- imagePullSecrets is added into all the Kubernetes resources
- Container images are updated
- Secret environment variables (for example api-token, signing-keys) now reference Kubernetes secrets
- Environment variables for RabbitMQ and MongoDB URIs change
- Environment variables for VM, OUTPUT and NOMAD service uri now reference <domain\_name>:<service\_port>
- Annotations from VM, OUTPUT and NOMAD service resources are deleted
- Github checksum is added as annotation
- Secret and annotation for distributor-\* deployments are deleted
- Upstream chart postgresql is updated
- Upstream charts will be recreated (delete and create):
  - Prometheus (circleci-server-kube-state-metrics, node-exporter, prometheus-server)
  - MongoDB
  - RabbitMQ
  - Redis (redis-master, redis-slave)

## 5. Upgrading CircleCI Server 3.x

Once your helm-value file is verified, run the following commands to upgrade the CircleCI server to v4.x.

Our helm registry is stored in an azure private registry. You will be provided a username and token to access the registry.

```
USER_NAME=<token-username>
PASSWORD=<token>
namespace=<your-install-namespace>
helm upgrade circleci-server -n <namespace> -f <path>/kots-exporter/output/helm-values.yaml
oci://cciserver.azurecr.io/circleci-server --version 4.0.0 --force --username $USER_NAME --password $PASSWORD
```

## 6. Check upgrade status

Run the following command to check all pods are up and running:

```
kubectl -n <namespace> get pods
```

## 7. Update DNS setting

Server 4.x migration is a destructive change for your DNS configuration. Server 4.x replaces the need for 4 load balancers and 5 DNS records with a single `load-balancer/external-ip`` service, named `circleci-proxy` or `circleci-proxy-acm`. This load balancer only needs to be routed via 2 DNS records, `<your-domain>` and `app.<your-domain>`. Separate domains for `vm-service`, `output-processor` and `nomad` are no longer required. Retrieve the external IP/Loadbalancer and update your DNS records accordingly.

```
kubectl -n <namespace> get svc circleci-proxy

# AWS Provider: XXXXX.elb.XXXXX.amazonaws.com
# GCP Provider: XXX.XXX.XXX.XXX
```

The following Kubernetes service objects are renamed/changed:

- `circleci-server-traefik` (LoadBalancer) → `kong` (ClusterIP)
- `nomad-server-external` (LoadBalancer) → `nomad-server` (ClusterIP)
- `output-processor` (LoadBalancer) → `output-processor` (ClusterIP)
- `vm-service` (LoadBalancer) → `vm-service` (ClusterIP)

The following Kubernetes service object is added:

- `circleci-proxy` or `circleci-proxy-acm` (LoadBalancer)

## 8. Execute Nomad Terraform

Execute the [Nomad Terraform](#) to re-create nomad client where `server_endpoint` value is set to be `<domain>:4647`. You can follow the steps mentioned [here](#). Update the helm value file with generated Certificates and Keys (base64 encoded) for Nomad Server-Client communication.

## 9. Validate your migration to server v4.x

Re-run [realitycheck](#) on your new server 4.x environment by pushing a fresh commit.

## 10. Update your team

Once you have successfully run [realitycheck](#), notify your team about the upgrade.